

Assignment 3

Experiments on Network slicing, SDN, ComNet simulator report

Jaishana Bindhu Priya (21bcs045)

Pawar Parth Sanjay (21bcs083)

R Eswar Naik (21bec035)

D.Rishikesh(21bec013)

R Hari prasad (21bec018)

ComNetsEmu Setup and Network Slicing:

We need to clone the repository from the GitLab public repository -

<https://git.comnets.net/public-repo/comnetsemu.git>

Once we have cloned the repository we can see the structure as shown below

Name	Last commit	Last update
└ .github/workflows	Port provision shell scripts to Ansible playbooks	1 year ago
└ app	Update flowvisor related scripts in multi_tenant_sdn_slic...	2 years ago
└ bin	Use sphinx for better API documentation	4 years ago
└ comnetsemu	Bump up to 0.3.1	1 year ago
└ doc	Update docs.yml	2 years ago
└ examples	Merge docs.yml to ci.yml	2 years ago
└ patch/mininet	Port provision shell scripts to Ansible playbooks	1 year ago
└ test_containers	Improve tools/scripts for setting up the test VM	2 years ago
└ util	Bump up to 0.3.1	1 year ago
◆ .gitattributes	<code_base>; Cleanup codebase, use a separate modu...	4 years ago
◆ .gitignore	Merge docs.yml to ci.yml	2 years ago
↳ .pylint	Improve code quality using Pylint.	4 years ago
↳ .shellcheckrc	<fix>; Fix warnings of shellcheck.	3 years ago
↳ CHANGELOG.md	Bump up to 0.3.1	1 year ago
↳ CONTRIBUTORS	Update CONTRIBUTORS.	3 years ago
◆ LICENSE	Revise README and tests, add LICENSE	4 years ago
↳ Makefile	Merge docs.yml to ci.yml	2 years ago
↳ README.md	Add reference for the book	6 months ago
↳ Vagrantfile	Bump up to 0.3.1	1 year ago
⚡ setup.py	Update documentation and dependencies to build Spil...	2 years ago

The examples that we will be looking into are present in the location -

comnetsemu/app/realizing_network_slicing

Now We will be looking into the working of different functions and APIs created to make this objective of Virtualizing the Network.

There are some basic requirements to run ComNetsEmu

- VirtualBox (or any other Virtual Machine)
- Vagrant
- Ubuntu OS

We need to run the following commands to start the virtualbox along with the ComNetsEmu :

```
vagrant up comnetsemu
```

- This will start up the virtualbox with username and password as comnetsemu.
- Once this runs correctly a confirmation message in green font appears stating that the virtualbox has been completed successfully. Next we run the following command.

```

Mar 28 5:53 PM
vagrant up commnetsemu

> cd commnetsemu
> vagrant up commnetsemu
Bringing machine 'commnetsemu' up with 'virtualbox' provider...
==> commnetsemu: Box 'bento/ubuntu-20.04' could not be found. Attempting to find and install...
    commnetsemu: Box Provider: virtualbox
    commnetsemu: Box Version: >= 0
==> commnetsemu: Loading metadata for box 'bento/ubuntu-20.04'
    commnetsemu: URL: https://vagrantcloud.com/api/v2/vagrant/bento/ubuntu-20.04
==> commnetsemu: Adding box 'bento/ubuntu-20.04' (v202401.31.0) for provider: virtualbox (amd64)
    commnetsemu: Downloading: https://vagrantcloud.com/bento/boxes/ubuntu-20.04/versions/202401.31.0/providers/virtualbox/amd64/vagrant.box
==> commnetsemu: Successfully added box 'bento/ubuntu-20.04' (v202401.31.0) for 'virtualbox (amd64)'
==> commnetsemu: Importing base box 'bento/ubuntu-20.04'...
==> commnetsemu: Matching MAC address for NAT networking...
==> commnetsemu: Checking if box 'bento/ubuntu-20.04' version '202401.31.0' is up to date...
==> commnetsemu: Setting the name of the VM: ubuntu-20.04-commnetsemu
Vagrant is currently configured to create VirtualBox synced folders with
the 'SharedFoldersEnableSymlinksCreate' option enabled. If the Vagrant
guest is not trusted, you may want to disable this option. For more
information on this option, please refer to the VirtualBox manual:
    https://www.virtualbox.org/manual/ch04.html#sharedfolders

This option can be disabled globally with an environment variable:
    VAGRANT_DISABLE_VBOXSYMLINKCREATE=1

or on a per folder basis within the Vagrantfile:
    config.vm.synced_folder '/host/path', '/guest/path', SharedFoldersEnableSymlinksCreate: false
==> commnetsemu: Clearing any previously set network interfaces...
==> commnetsemu: Preparing network interfaces based on configuration...
    commnetsemu: Adapter 1: nat
==> commnetsemu: Forwarding ports...
    commnetsemu: 8888 (guest) => 8888 (host) (adapter 1)
    commnetsemu: 8082 (guest) => 8082 (host) (adapter 1)
    commnetsemu: 8083 (guest) => 8083 (host) (adapter 1)
    commnetsemu: 8084 (guest) => 8084 (host) (adapter 1)
    commnetsemu: 22 (guest) => 2222 (host) (adapter 1)
==> commnetsemu: Running 'pre-boot' VM customizations...
==> commnetsemu: Booting VM...
==> commnetsemu: Waiting for machine to boot. This may take a few minutes...
    commnetsemu: SSH address: 127.0.0.1:2222
    commnetsemu: SSH username: vagrant
    commnetsemu: SSH auth method: private key
    commnetsemu: Warning: Connection reset. Retrying...
    commnetsemu: Warning: Remote connection disconnect. Retrying...
    commnetsemu:
    commnetsemu: Vagrant insecure key detected. Vagrant will automatically replace
    commnetsemu: this with a newly generated keypair for better security.
    commnetsemu:
    commnetsemu: Inserting generated public key within guest...
    commnetsemu: Removing insecure key from the guest if it's present...

```

```

commnetsemu:
==> commnetsemu: Running provisioner: shell...
    commnetsemu: Running: inline script
    commnetsemu: Reading package lists...
    commnetsemu: Building dependency tree...
    commnetsemu: Reading state information...
    commnetsemu: Reading package lists...
    commnetsemu: Building dependency tree...
    commnetsemu: Reading state information...
    commnetsemu: 0 upgraded, 0 newly installed, 0 to remove and 3 not upgraded.
==> commnetsemu: Running provisioner: shell...
    commnetsemu: Running: /tmp/vagrant-shell20240328-13111-gr5bou.sh
    commnetsemu:
    commnetsemu:
    commnetsemu:
    commnetsemu:
    commnetsemu:
    commnetsemu:
    commnetsemu:
    commnetsemu: 
    commnetsemu:
    commnetsemu: Machine 'commnetsemu' has a post `vagrant up` message. This is a message
    commnetsemu: from the creator of the Vagrantfile, and not from Vagrant itself:
    commnetsemu:
    commnetsemu:
    commnetsemu: The VM is up! Run "vagrant ssh commnetsemu" to ssh into the running VM.
    commnetsemu:
    commnetsemu: IMPORTANT! For all ComNetSEmu users and developers:
    commnetsemu:
    commnetsemu: When a new version is released (https://git.comnets.net/public-repo/commnetsemu/-/tags),
    commnetsemu: please run the upgrade process described [here](<a href="https://git.comnets.net/public-repo/commnetsemu#upgrade-commnetsemu-and-dependencies">here</a>).
    commnetsemu: New features, fixes and other improvements require **manually** running the upgrade script.
    commnetsemu: But the script will automatically check and perform the upgrade, and if you have a good internet connection,
    commnetsemu: it will not take much time.
    commnetsemu:

```

vagrant ssh commnetsemu

- This loads up the virtual machine which opens a ssh (command line) interface so that we can run our commands in the VM.

- We can run `cd comnetsemu` and then try to run to check if our installation was successful by executing the “docker-in-docker” project.

```
> vagrant ssh comnetsemu
Welcome to Ubuntu 20.04.6 LTS (GNU/Linux 5.4.0-170-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/pro

 System information as of Thu 28 Mar 2024 12:34:37 PM UTC

 System load:  0.36          Processes:           162
 Usage of /:   18.2% of 30.34GB  Users logged in:      0
 Memory usage: 14%          IPv4 address for docker0: 172.17
 Swap usage:   0%          IPv4 address for eth0:    10.0.2

This system is built by the Bento project by Chef Software
More information can be found at https://github.com/chef/bento
/usr/bin/xauth:  file /home/vagrant/.Xauthority does not exist
vagrant@comnetsemu:~$ 
```

`./install.sh -u`

- This will install the ComNetsemu Emulator.
- Once the task is completed, it confirms to delete and close all the images that are running currently.

```
TASK [Build all test containers for ComNetsEmu (required for unit tests and some built-in examples)] ****
[WARNING]: Skipping plugin (/usr/lib/python3/dist-packages/ansible/plugins/filter/core.py) as it seems to be invalid: cannot import name 'environmentfilter'
(/usr/local/lib/python3.8/dist-packages/jinja2/filters.py)
[WARNING]: Skipping plugin (/usr/lib/python3/dist-packages/ansible/plugins/filter/mathstuff.py) as it seems to be invalid: cannot import name 'environment'
(/usr/local/lib/python3.8/dist-packages/jinja2/filters.py)
changed: [127.0.0.1]

TASK [Install packages required for the development of ComNetsEmu] ****
ok: [127.0.0.1]

TASK [Install packages required to build HTML documentation of ComNetsEmu] ****
[WARNING]: Skipping plugin (/usr/lib/python3/dist-packages/ansible/plugins/filter/core.py) as it seems to be invalid: cannot import name 'environmentfilter'
(/usr/local/lib/python3.8/dist-packages/jinja2/filters.py)
[WARNING]: Skipping plugin (/usr/lib/python3/dist-packages/ansible/plugins/filter/mathstuff.py) as it seems to be invalid: cannot import name 'environment'
(/usr/local/lib/python3.8/dist-packages/jinja2/filters.py)
changed: [127.0.0.1]

TASK [Start and enable the Docker service] ****
ok: [127.0.0.1]

PLAY RECAP ****
127.0.0.1 : ok=14  changed=7  unreachable=0  failed=0  skipped=1  rescued=0  ignored=0

vagrant@comnetsemu:~/comnetsemu/util$ 
```

`sudo make test`

- This will try to run unit tests to check the installation of all the needed components.
- If the final print statement says “Ran 5 tests in 86.6s” and “OK”. This means that all the preliminary tests have been passed and we can carry on with the next step to try and execute the docker-in-docker project to test out and continue with our other objectives.

```
vagrant@comnetsemu:~/comnetsemu/util$ docker image ls
REPOSITORY      TAG      IMAGE ID      CREATED      SIZE
dev_test        latest    0bb965fe1e45  7 minutes ago  113MB
ubuntu          20.04    3cff1c6ff37e  5 weeks ago   72.8MB
vagrant@comnetsemu:~/comnetsemu/util$ cd ..
vagrant@comnetsemu:~/comnetsemu$ ls
app  bin  CHANGELOG.md  comnetsemu  CONTRIBUTORS  doc  examples  L1
vagrant@comnetsemu:~/comnetsemu$ sudo make test
Run all unit tests of ComNetsEmu Python package.
python3 ./comnetsemu/test/unit/runner.py -v
test_appcontainermanager_rest_api (test_net.TestVNFManager) ... ok
test_container_crud (test_net.TestVNFManager) ... ok
test_container_isolation (test_net.TestVNFManager) ... ok
test_ping (test_net.TestVNFManager) ... ok
test_appcontainer (test_node.TestComNetsEmuNode) ... ok
-----
Ran 5 tests in 20.112s
OK
```

make doc

- This compiles all the important functions and provides an HTML page to show us the required functions and APIs.
- This contains all the necessary information regarding the components of ComNetsEmu.

```

OK
vagrant@comnetsemu:~/comnetsemu$ make doc
Build documentation in HTML format
rm ./doc/api/*.rst
rm: cannot remove './doc/api/*.rst': No such file or directory
make: [Makefile:84: doc] Error 1 (ignored)
rm -r ./doc/build/
rm: cannot remove './doc/build/': No such file or directory
make: [Makefile:85: doc] Error 1 (ignored)
# Use sphinx-apidoc to generate API sources
cd ./doc/ && sphinx-apidoc -o api .. /comnetsemu -H 'Python API' --separate
Creating file api/comnetsemu.rst.
Creating file api/comnetsemu.clean.rst.
Creating file api/comnetsemu.cli.rst.
Creating file api/comnetsemu.exceptions.rst.
Creating file api/comnetsemu.net.rst.
Creating file api/comnetsemu.node.rst.
Creating file api/comnetsemu.overrides.rst.
Creating file api/comnetsemu.tool.rst.
Creating file api/comnetsemu.util.rst.
Creating file api/modules.rst.
# Build HTML files
cd ./doc/ && make html
make[1]: Entering directory '/home/vagrant/comnetsemu/doc'
Running Sphinx v7.1.2
WARNING: Invalid configuration value found: 'language = None'. Update your config
making output directory... done
myst v2.0.0: MdParserConfig(commonmark_only=False, gfm_only=False, enable_extensions=domains=None, fence_as_directive=set(), number_code_blocks=[], title_to_header=False, substitutions={}, linkify_fuzzy_links=True, dmath_allow_labels=True, dmath_allow_labels_process='math|mathjax_process|math|output_area', enable_checkboxes=False, suppress_warnings=True)
building [mo]: targets for 0 po files that are out of date
writing output...
building [html]: targets for 16 source files that are out of date
updating environment: [new config] 16 added, 0 changed, 0 removed
reading sources... [100%] vm_setup_issues
looking for now-outdated files... none found
pickling environment... done
checking consistency... done
preparing documents... done
copying assets... copying static files... done
copying extra files... done
done
writing output... [100%] vm_setup_issues
generating indices... genindex py-modindex done
writing additional pages... search done
dumping search index in English (code: en)... done
dumping object inventory... done
build succeeded, 1 warning.

The HTML pages are in build/html.
make[1]: Leaving directory '/home/vagrant/comnetsemu/doc'

```

ComNetsEmu

A virtual emulator/testbed designed for the book: Computing in Communication Networks: From Theory to Practice

Navigation

- [Python API](#)
- [Installation](#)
- [FAQ](#)
- [Reference Links](#)
- [Roadmap of ComNetsEmu](#)
- [Known VM Setup Issues](#)

Quick search

ComNetsEmu's Documentation!

A virtual emulator/testbed designed for the book: [Computing in Communication Networks: From Theory to Practice](#)

Check the [README](#) for project description.

This online page is mainly for ComNetsEmu API documentation.

- [Python API](#)
 - [comnetsemu package](#)
- [Installation](#)
 - [Option 1: Install in a Vagrant managed VM \(Highly Recommended\)](#)
 - [Option 2: Install on user's custom VM or directly on host OS](#)
 - [Option 3: Download the pre-built VM image](#)
 - [Post-Installation](#)
- [FAQ](#)

```

class comnetsemu.net.APPContainerManager(net: Mininet)
    Bases: object

    Manager for application containers (sibling containers) deployed on Mininet hosts.

    addContainer(name: str, dhost: str, dimage: str, dcmd: str,
    docker_args: dict | None = None, wait: bool = True) →
    APPContainer
        Create and run a new container inside a Mininet DockerHost.

        Parameters: • name (str) – Name of the container.
                      • dhost – Name of the host used for deployment.
                      • dimage (str) – Name of the docker image.
                      • dcmd (str) – Command to run after the creation.
                      • docker_args (dict) – All other keyword arguments supported by
                        Docker-py. e.g. CPU and memory related limitations. Some parameters
                        are overridden for APPContainerManager's functionalities.
                      • wait (bool) – Wait until the container has the running state if True.

        Check cls.docker_args_default.

        Returns: Added APPContainer instance or None if the creation process failed.
        Return type: APPContainer

    docker_args_default = {'detach': True, 'init': True,
    'labels': {'comnetsemu': 'dockercontainer'}, 'security_opt':
    ['seccomp:unconfined'], 'tty': True}

```

- The above image shows the component - **addContainer**, along with the parameters it accepts and helps us understand the working of the components.
- These components are built on the basis of MiniNet (another Virtualization software)

Now we will discuss the code and structure of dockerindocker.py in the examples folder in comnetsemu folder.

```
vagrant@comnetsemu:~/comnetsemu/examples$ ls -a
.  dockerhostcli.py      dockerhost.py      dpdk       flowvisor      network_measurement.py  README.md  service_migration  vpn
..  dockerhost_manage_appcontainer.py  dockerindocker.py  echo_server  mininet_demystify  ovx_test.py   run.sh    tun_ebpf.py
vagrant@comnetsemu:~/comnetsemu/examples$ 
```

- We will now look into the code and try to understand what the important functions in dockerindocker.py perform.

- Below is the structure of the network that we will be trying to replicate.
(Chain Topology)

```
GNU nano 4.8
#!/usr/bin/python3

"""
About: Basic example of running Docker container inside Docker container Host

Topo: Chain topology   h1      h2      h3          hn
           |        |        |
           s1 --- s2 --- s3 --- ... sn

Tests:
- Ping test between internal container between head (h1) and tail (hn).
- Resource limitation test for head and tail hosts: The CPU and memory
  limitation of Dockerhost are configured to (50/n)% and 10MB. The internal
  containers try to use 100% and 300MB RAM by execute stree-nng. The actual
  resource usages are monitored and reported.
"""


```

- Now in the main function we will initialize some important variables like Number of hosts (host_num)
- Then we call the function testDockerInDocker function and pass the required parameters.

```
if __name__ == "__main__":
    setLogLevel("info")
    host_num = 3
    parser = argparse.ArgumentParser(
        description="Basic example for Docker-in-Docker setup."
    )
    parser.add_argument(
        "--host_num", default=3, type=int, help="Number of hosts in the chain topology"
    )
    args = parser.parse_args()
    print("!!! The number of hosts in the chain: " + str(args.host_num))
    testDockerInDocker(args.host_num)
    
```

- Now we add the hosts to the network.
- We add the switches as well to the network and also provide the bandwidth and delay for each switch.
- Once this is completed we Start the Network.

```

def testDockerInDocker(n):

    net = Containernet(controller=Controller, link=TCLink)
    mgr = VNFManager(net)

    info("*** Adding controller\n")
    net.addController("c0")

    info("*** Adding Docker hosts and switches in a chain topology\n")
    last_sw = None
    hosts = list()
    # Connect hosts
    for i in range(n):
        # Unlimited access to CPU(s) cycles in CPU_SETS
        host = net.addDockerHost(
            "h%s" % (i + 1),
            dimage="dev_test",
            ip="10.0.0.%s" % (i + 1),
            docker_args={"cpuset_cpus": "0"},
        )
        hosts.append(host)
        switch = net.addSwitch("s%s" % (i + 1))
        net.addLink(switch, host, bw=10, delay="10ms")
        if last_sw:
            # Connect switches
            net.addLink(switch, last_sw, bw=10, delay="10ms")
        last_sw = switch

    info("*** Starting network\n")
    net.start()

    info(
        "*** Run a simple ping test between two internal containers deployed on h1 and h%s\n"
        % n
    )
    head = mgr.addContainer("head", "h1", "dev_test", "/bin/bash", docker_args={})
    tail = mgr.addContainer(
        "tail", "h%s" % n, "dev_test", "ping -c 3 10.0.0.1", docker_args={}
    )

```

- We also provide various parameters to the addDockerHost function,
- **dimage** : this parameter takes in the docker image needed to assign the host.
- **Docker_args** : this parameter some arguments like cpuset_cpus to 0, which means in which cpu the code should be executed.

- **cpuset_cpus (str)** – CPUs in which to allow execution (0-3, 0, 1).
- Now we try to ping the internal containers h1 and hn where n is the number of hosts, here it is 3.

```

info(
    "*** Run a simple ping test between two internal containers deployed on h1 and h%s\n"
    % n
)
head = mgr.addContainer("head", "h1", "dev_test", "/bin/bash", docker_args={})
tail = mgr.addContainer(
    "tail", "h%s" % n, "dev_test", "ping -c 3 10.0.0.1", docker_args={}
)

info("*** Tail start ping head, wait for 5s...")
time.sleep(5)
info("\nThe ping result of tail to head: \n")
print(tail.dins.logs().decode("utf-8"))
mgr.removeContainer(head.name)
mgr.removeContainer(tail.name)

time.sleep(3)

```

- We use the addConatiner to assign head and tail for the first and last host.
- As we can see in the image given below the function tries to a new docker within the already running docker container.
- This tries to make a copy named host1 (h1) and host3 (h3) which run on the docker image “dev_test”.
- They run on “/bin/bash” a shell and run “ping -c 3 10.0.0.1” respectively.
- This basically makes the head host ping the tail host.

comnetsemu.net module

About: ComNetsEmu Network Module.

```
class comnetsemu.net.APPContainerManager(net: Mininet)
    Bases: object

    Manager for application containers (sibling containers) deployed on Mininet hosts.

    addContainer(name: str, dhost: str, dimage: str, dcmd: str,
    docker_args: dict | None = None, wait: bool = True) →
    APPContainer
        Create and run a new container inside a Mininet DockerHost.

        Parameters: • name (str) – Name of the container.
                    • dhost – Name of the host used for deployment.
                    • dimage (str) – Name of the docker image.
                    • dcmd (str) – Command to run after the creation.
                    • docker_args (dict) – All other keyword arguments supported by
                      Docker-py. e.g. CPU and memory related limitations. Some parameters
                      are overridden for APPContainerManager's functionalities.
                    • wait (bool) – Wait until the container has the running state if True.

        Check cls.docker_args_default.

        Returns: Added APPContainer instance or None if the creation process failed.
        Return type: APPContainer

    docker_args_default = {'detach': True, 'init': True,
    'labels': {'comnetsemu': 'dockercontainer'}, 'security_opt':
    ['seccomp:unconfined'], 'tty': True}
```

- Next we need to set the CPU limitations for each of the hosts (docker containers).

```
info(
    "Update CPU limitation of all Docker hosts (h1-h%s) to %.2f %% \n"
    % (n, 50.0 / n)
)
info("Update Memory limitation of all Docker hosts (h1-h%s) to 10MB\n" % n)
info(
    "Deploy the stress app (100 % CPU and 300M memory) inside all Docker hosts to test the CPU"
)
containers = list()
# Mark: CPU percent = cpu_quota / cpu_period. The default CPU period is
# 100ms = 100000 us
for i, h in enumerate(hosts):
    h.dins.update(cpu_quota=int(50000 / n))
    h.dins.update(mem_limit=10 * (1024 ** 2)) # in bytes
    c = mgr.addContainer(
        "stress_app_%s" % (i + 1),
        h.name,
        "dev_test",
        "stress-ng -c 1 -m 1 --vm-bytes 300M",
        docker_args={}
    )
    containers.append(c)

info(
    "Start monitoring resource usage of internal containers"
    "with default sample count: 3 and sample period: 1 sec\n"
)
```

- As seen in the above image, we create a list named containers and then append c into it which basically uses addContainer() and updates the

maximum CPU usage based on CPU percent = `cpu_quota / cpu_period`
(which are defined in the below image).

- `cpu_percent (int)` – Usable percentage of the available CPUs (Windows only).
 - `cpu_period (int)` – The length of a CPU period in microseconds.
 - `cpu_quota (int)` – Microseconds of CPU time that the container can get in a CPU period.
-
- Now we try to set the `cpu_quota` and memory limitation. We also assign these to the `addContainer()` and add it to the container list.
 - We use a tool named `stress-ng` “`stress-ng -c 1 -m 1 -vm-bytes 300M`”. This means that we assign 1 core, 100% of the CPU and 300MB of memory or RAM.
 - Next we make use of a method called `monResourceStats` that basically monitors the Resources used by a container.

```
for c in containers:
    usages = mgr.monResourceStats(c.name)
    if usages:
        print(
            "The average CPU and Memory usage of container:{} is {:.2f}%, {:.2f}MB".format(
                c.name,
                (sum(u[0] for u in usages) / len(usages)),
                (sum(u[1] for u in usages) / len(usages)),
            )
        )
    else:
        print("[ERROR] Failed to get resource usages from manager")

for c in containers:
    mgr.removeContainer(c.name)

info("*** Stopping network\n")
net.stop()
mgr.stop()
```

- Here we iterate over all the containers and print out the usages of each container using the `monResourceStats()` function and print the CPU and Memory Usages.
- Once these are done we remove the containers and stop the network emulation.
- Now we run the example in the virtual machine in the example directory.

```
sudo python3 dockerindocker.py
```

```

vagrant@comnetsemu:~/comnetsemu/examples$ sudo python3 dockerindocker.py
*** The number of hosts in the chain: 3
*** Adding controller
*** Adding Docker hosts and switches in a chain topology
(10.00Mbit 10ms delay) (10.00Mbit 10ms delay) (10.00Mbit 10ms delay) (10.00Mbit 10ms delay) (10.00Mbit 10ms delay)
Mbit 10ms delay) (10.00Mbit 10ms delay) *** Starting network
*** Configuring hosts
h1 h2 h3
*** Starting controller
c0
*** Starting 3 switches
s1 (10.00Mbit 10ms delay) (10.00Mbit 10ms delay) s2 (10.00Mbit 10ms delay) (10.00Mbit 10ms delay) (10.00Mbit
delay) (10.00Mbit 10ms delay) (10.00Mbit 10ms delay) (10.00Mbit 10ms delay) (10.00Mbit 10ms delay)
*** Run a simple ping test between two internal containers deployed on h1 and h3
*** Tail start ping head, wait for 5s...
The ping result of tail to head:
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
64 bytes from 10.0.0.1: icmp_seq=1 ttl=64 time=196 ms
64 bytes from 10.0.0.1: icmp_seq=2 ttl=64 time=101 ms
64 bytes from 10.0.0.1: icmp_seq=3 ttl=64 time=355 ms

--- 10.0.0.1 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2005ms
rtt min/avg/max/mdev = 100.618/217.278/355.232/105.030 ms

*** Run CPU resource limitation test. Update CPU limitation of all Docker hosts (h1-h3) to 16.67 %
Update Memory limitation of all Docker hosts (h1-h3) to 10MB
Deploy the stress app (100 % CPU and 300M memory) inside all Docker hosts to test the CPU/Memory limitation
Start monitoring resource usage of internal containers with default sample count: 3 and sample period: 1 sec
    The average CPU and Memory usage of container:stress_app_1 is 17.75%, 9.03MB
    The average CPU and Memory usage of container:stress_app_2 is 17.44%, 8.94MB
    The average CPU and Memory usage of container:stress_app_3 is 17.53%, 8.91MB
*** Stopping network
*** Stopping 1 controllers
c0
*** Stopping 5 links
.....
*** Stopping 3 switches
s1 s2 s3
*** Stopping 3 hosts
h1 h2 h3
*** Done

```

- The above is the results obtained after executing the file.
- Here we can make the following observations :
- First packet takes more time (nearly 1815ms) due to initialization of the app.
- We also limited the maximum CPU usages of each container to sum upto a maximum of 50%.
- In the monitoring function we can observe that we use a maximum of 16% of CPU and 10MB of RAM.
-
- We can see that all the containers are working based on the constraints that we provideR

In another example, dockerhost_manage_appcontainer.py

- Here we try to manage the application container from one Docker Host Instance rather than directly from the manager.
- Here docker0 bridges DockerHosts and AppContainerManager. We also make use of REST APIs to send requests to create and delete an APP container on the master node h1.

```
GNU nano 4.8
#!/usr/bin/env python3
# -*- coding: utf-8 -*-
# vim:fenc=utf-8

"""

About: This example shows how to manage application containers from one
DockerHost instance instead of calling functions directly on manager
object.
This can be useful if you want to let one DockerHost node in the network
to be the master node (like the master node in K8s) in to manage
application containers on other worker nodes.

In this example, h1 is the master node. h2 wants to create a APP
container srv2_1 on it by sending ping packets to h1.
When h1 receive the ping packet from h2, it uses REST API to request a
creation of srv2_1 to the APPContainerManager's HTTP server.
h1 also uses the DELETE request to remove the srv1_1 (created by the
native API) on itself.
Since all DockerHosts are connected to the docker0 bridge by default,
docker0 is used as the bridge to connect DockerHosts and the
AppContainerManager.

For the REST API implementation, check commetsemu/net.py for details.

"""


```

- Now we will go through the code to explain the working of function AppConatinerManager()

```

if __name__ == "__main__":
    # Only used for auto-testing.
    AUTOTEST_MODE = os.environ.get("COMNETSEMU_AUTOTEST_MODE", 0)

    setLogLevel("info")

    net = Containernet(controller=Controller, link=TCLink, xterms=False)
    mgr = VNFManager(net)

    info("*** Add controller\n")
    net.addController("c0")

    info("*** Creating hosts\n")
    h1 = net.addDockerHost(
        "h1", dimage="dev_test", ip="10.0.0.1", docker_args={"hostname": "h1"}
    )
    h2 = net.addDockerHost(
        "h2", dimage="dev_test", ip="10.0.0.2", docker_args={"hostname": "h2"}
    )

    info("*** Adding switch and links\n")
    switch1 = net.addSwitch("s1")
    switch2 = net.addSwitch("s2")
    net.addLink(switch1, h1, bw=10, delay="10ms")
    net.addLink(switch1, switch2, bw=10, delay="10ms")
    net.addLink(switch2, h2, bw=10, delay="10ms")

    info("\n*** Starting network\n")
    net.start()

    ip_route = pyroute2.IPRoute()
    mgr_api_ip = ip_route.get_addr(label="docker0")[0].get_attr("IFA_ADDRESS")
    print("*** Deploy srv1_1 on h1 with native API.")
    mgr.addContainer("srv1_1", "h1", "dev_test", "bash", docker_args={})
    print_deployed_containers(mgr, ["h1", "h2"])

```

- As we can see above, It is similar to the previous example.
- The main change here is that :
- We deploy the master node directly.
- But we also need to manually run the HTTP server by providing the port and IP address.

```

print(" *** Run REST API server thread. It listens on docker0 interface")
mgr.runRESTServerThread(ip=mgr_api_ip, port=8000)
time.sleep(1)

```

- We can use the REST API to add and remove the APP from h1.

- We are using TCP dump to print out the result after the ping to the master node h1.

```

request_url = f"{mgr_api_ip}:8000/containers"
print("**** h1 waits for one Ping message from h1.")
h2.cmd("ping 10.0.0.1 &")
ret = h1.cmd("tcpdump -i h1-eth0 -c 1 icmp")
print(f"h1: output of tcpdump:\n {ret}")
print("**** h1 request to create a container named srv2_1 on h2 via REST API.")
h1.cmd(
    'curl -X POST --data \'{"name": "srv2_1", "dhost": "h2", "dimage": "dev_test", "dcmd": "bash", "docker_args": {}}\'' + request_url
)
time.sleep(1)
print_deployed_containers(mgr, ["h1", "h2"])

print("**** h1 request to delete a container named srv1_1 on itself via REST API.")
request_url += "/srv1_1"
h1.cmd("curl -X DELETE " + request_url)
time.sleep(1)
print_deployed_containers(mgr, ["h1", "h2"])

```

We run the file by using the following command

```
sudo python3 dockerhost_manage_appcontainer.py
```

```

vagrant@comnetsemu:~/comnetsemu/examples$ sudo python3 dockerhost_manage_appcontainer.py
*** Add controller
*** Creating hosts
*** Adding switch and links
(10.00Mbit 10ms delay) (10.00Mbit 10ms delay) (10.00Mbit 10ms delay) (10.00Mbit 10ms delay)
*** Starting network
*** Configuring hosts
h1 h2
*** Starting controller
c0
*** Starting 2 switches
s1 (10.00Mbit 10ms delay) (10.00Mbit 10ms delay) s2 (10.00Mbit 10ms delay) (10.00Mbit 10ms delay)
*** Deploy srv1_1 on h1 with native API.
*** Current deployed APP containers:
DockerHost: h1, APP containers on it: srv1_1
DockerHost: h2, APP containers on it:
*** Run REST API server thread. It listens on docker0 interface
Start REST API server on address: 172.17.0.1:8000.
**** h1 waits for one Ping message from h1.
h1: output of tcpdump:
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on h1-eth0, link-type EN10MB (Ethernet), capture size 262144 bytes
13:02:32.108128 IP 10.0.0.2 > 10.0.0.1: ICMP echo request, id 8, seq 1, length 64
1 packet captured
26 packets received by filter
0 packets dropped by kernel

**** h1 request to create a container named srv2_1 on h2 via REST API.
172.17.0.2 - - [28/Mar/2024 13:02:44] "POST /containers HTTP/1.1" 200 -
*** Current deployed APP containers:
DockerHost: h1, APP containers on it: srv1_1
DockerHost: h2, APP containers on it: srv2_1
**** h1 request to delete a container named srv1_1 on itself via REST API.
172.17.0.2 - - [28/Mar/2024 13:02:45] "DELETE /containers/srv1_1 HTTP/1.1" 200 -
*** Current deployed APP containers:
DockerHost: h1, APP containers on it:
DockerHost: h2, APP containers on it: srv2_1
*** Starting CLI:

```

- As we can see in the above image,
- Once the ping is received on the master node h1, it creates an APP named “srv2_1” and we also remove the APP named “srv1_1”.
- We can exit the docker once we type “exit” in the miniNet CLI.

```
mininet> exit
*** Stopping 1 controllers
c0
*** Stopping 3 links
...
*** Stopping 2 switches
s1 s2
*** Stopping 2 hosts
h1 h2
*** Done
Stop 1 containers in the App container queue: srv2_1
vagrant@comnetsemu:~/comnetsemu/examples$ 
```

Open vSwitch Exercise 1: Create a bridge in OpenvSwitch

The focus of the exercise is the creation and configuration of a network bridge using Open vSwitch (OVS).

Environment:

- Virtual Machine: The exercise is conducted within a "comnetsemu" virtual machine, likely running on VirtualBox or a similar platform. This VM is specifically tailored for network emulation and experimentation.
- ComNetSEmu Framework: ComNetSEmu provides the underlying tools and environment for conducting network simulations.
- Mininet: A popular network emulator built upon Linux namespaces and utilities, Mininet allows users to create realistic network topologies within a single system.
- OpenFlow Controller: This exercise involves an OpenFlow controller, likely running on the same machine (localhost - 127.0.0.1), which is responsible for managing and controlling the flow of network traffic through the OVS bridge.

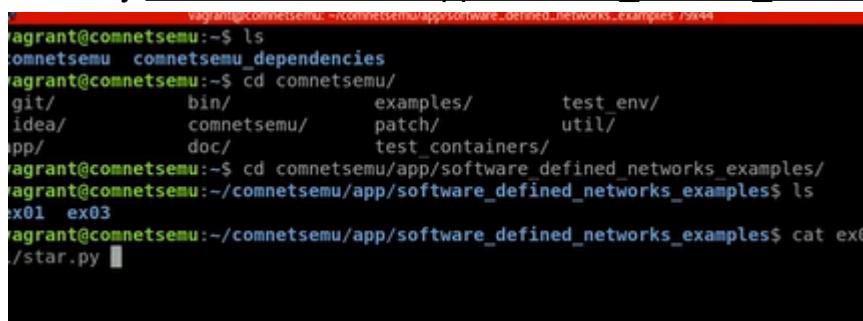
Exercise Steps:

Terminal Access & VM Login:

- Access a terminal on the host machine.
- Log into the ComNetSEmu virtual machine using the command: vagrant ssh comnetsemu.

Navigation to Exercise Directory:

- Within the VM, navigate to the exercise directory:cd~/comnetsemu/app/software_defined_networks_examples/



```
vagrant@comnetsemu:~$ ls
comnetsemu  comnetsemu_dependencies
vagrant@comnetsemu:~$ cd comnetsemu/
git/          bin/        examples/      test_env/
idea/         comnetsemu/  patch/       util/
ipp/          doc/        test_containers/
vagrant@comnetsemu:~$ cd comnetsemu/app/software_defined_networks_examples/
vagrant@comnetsemu:~/comnetsemu/app/software_defined_networks_examples$ ls
ex01  ex03
vagrant@comnetsemu:~/comnetsemu/app/software_defined_networks_examples$ cat ex01/star.py
```

Python Script Execution:

- Execute the Python script to initiate the exercise scenario: sudo python3 ex01/star.py

```

"""
from mininet.net import Mininet
from mininet.node import RemoteController
from mininet.cli import CLI
from mininet.log import setLogLevel, info
from nullswitch import NullSwitch

def emptyNet():

    "Create an empty network and add nodes to it."

    net = Mininet( controller=RemoteController, switch=NullSwitch )

    info( "*** Adding controller\n" )
    net.addController( 'c0' )

    info( "*** Adding hosts\n" )
    h1 = net.addHost( 'h1' )
    h2 = net.addHost( 'h2' )
    h3 = net.addHost( 'h3' )

    info( "*** Adding switch\n" )
    s1 = net.addSwitch( 's1' )

    info( "*** Creating links\n" )
    net.addLink( h1, s1 )
    net.addLink( h2, s1 )
    net.addLink( h3, s1 )

    info( "*** Starting network\n" )
    net.start()

    info( "*** Running CLI\n" )
    CLI( net )

    info( "*** Stopping network" )
    net.stop()

if __name__ == '__main__':
    setLogLevel( 'info' )
    emptyNet()

```

- This script likely sets up the initial Mininet network topology.

```
vagrant@comnetsemu:~/comnetsemu/app/software_defined_networks_examples$ sudo python3 ex01/star.py
*** Adding controller
Unable to contact the remote controller at 127.0.0.1:6653
Unable to contact the remote controller at 127.0.0.1:6633
Setting remote controller to 127.0.0.1:6653
*** Adding hosts
*** Adding switch
*** Creating links
*** Starting network
*** Configuring hosts
h1 h2 h3
*** Starting controller
c0
*** Starting 1 switches
s1
*** Running CLI
*** Starting CLI:
mininet> █
```

Mininet Shell Interaction:

- The user is now within the Mininet CLI.
- pingall: This command verifies network connectivity by sending ICMP (ping) requests between nodes in the emulated network.

```
mininet> pingall
*** Ping: testing ping reachability
h1 -> X X
h2 -> X X
h3 -> X X
*** Results: 100% dropped (0/6 received)
mininet> █
```

- s1 xterm &: Opens an xterm (terminal window) connected to node "s1," enabling direct interaction with the node.

```
mininet> pingall
*** Ping: testing ping reachability
h1 -> X X
h2 -> X X
h3 -> X X
*** Results: 100% dropped (0/6 received)
mininet> s1 xterm &
```

Bridge Creation:

- The central task is creating an OVS bridge named "s1", acting as a virtual switch connecting multiple network interfaces.

- OVS commands used for bridge creation and interface addition:`ovs-vsctl add-br s1` (Adds bridge "s1")

```

        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 52:54:00:30:ab:11 brd ff:ff:ff:ff:ff:ff
    inet 192.168.121.219/24 brd 192.168.121.255 scope global dynamic eth0
        valid_lft 3342sec preferred_lft 3342sec
    inet6 fe80::5054:ff:fe30:ab11/64 scope link
        valid_lft forever preferred_lft forever
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
    link/ether 02:42:68:c5:40:b2 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
4: s1-eth1@if2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 4e:0e:3d:ea:22:1f brd ff:ff:ff:ff:ff:ff link-netnsid 0
5: s1-eth2@if2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether a2:e2:2e:60:0b:06 brd ff:ff:ff:ff:ff:ff link-netnsid 1
6: s1-eth3@if2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether e2:89:a2:67:66:a2 brd ff:ff:ff:ff:ff:ff link-netnsid 2
root@comnetsemu:~/comnetsemu/app/software_defined_networks_examples# ovs
unning CLI

```

- `ovs-vsctl add-port s1 s1-eth1` (Adds interface "s1-eth1" to the bridge)

```

root@comnetsemu:~/comnetsemu/app/software_defined_networks_examples (on comnetsemu)
        valid_lft forever preferred_lft forever
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
    link/ether 02:42:68:c5:40:b2 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
4: s1-eth1@if2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 4e:0e:3d:ea:22:1f brd ff:ff:ff:ff:ff:ff link-netnsid 0
5: s1-eth2@if2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether a2:e2:2e:60:0b:06 brd ff:ff:ff:ff:ff:ff link-netnsid 1
6: s1-eth3@if2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether e2:89:a2:67:66:a2 brd ff:ff:ff:ff:ff:ff link-netnsid 2
root@comnetsemu:~/comnetsemu/app/software_defined_networks_examples# ovs-vsctl s
how
138f8106-c225-4f78-aacc-f2d8724bb3ef
    ovs_version: "2.9.5"
root@comnetsemu:~/comnetsemu/app/software_defined_networks_examples# ovs-vsctl show
138f8106-c225-4f78-aacc-f2d8724bb3ef
    ovs_version: "2.9.5"

```

- ovs-vsctl add-port s1 s1-eth2 (Adds interface "s1-eth2" to the bridge)

```
ovs_version: 2.9.5
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# ovs-vsctl add-port s1 s1-eth1
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# ovs-vsctl add-port s1 s1-eth2
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# ovs-vsctl add-port s1 s1-eth3
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# ovs-vsctl show
438f8106-c225-4f78-aacc-f2d8724bb3ef
    Bridge "s1"
        Port "s1-eth2"
            Interface "s1-eth2"
        Port "s1-eth1"
            Interface "s1-eth1"
        Port "s1"
            Interface "s1"
                type: internal
        Port "s1-eth3"
            Interface "s1-eth3"
    ovs_version: "2.9.5"
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples#
```

- ovs-vsctl add-port s1 s1-eth3 (Adds interface "s1-eth3" to the bridge)

```
ovs_version: 2.9.5
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# ovs-vsctl add-port s1 s1-eth1
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# ovs-vsctl add-port s1 s1-eth2
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# ovs-vsctl add-port s1 s1-eth3
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# ovs-vsctl show
438f8106-c225-4f78-aacc-f2d8724bb3ef
    Bridge "s1"
        Port "s1-eth2"
            Interface "s1-eth2"
        Port "s1-eth1"
            Interface "s1-eth1"
        Port "s1"
            Interface "s1"
                type: internal
        Port "s1-eth3"
            Interface "s1-eth3"
    ovs_version: "2.9.5"
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples#
```

Verification:

After bridge creation and interface connections, pingall is used again within the Mininet shell to confirm successful communication between nodes via the new bridge.

```
root@comnetsemu:~/comnetsemu/app/software_defined_networks_examples# ovs-vsctl set-controller s1 tcp:127.0.0.1:6653
root@comnetsemu:~/comnetsemu/app/software_defined_networks_examples# ovs-vsctl show
438f8106-c225-4f78-aacc-f2d8724bb3ef
  Bridge "s1"
    Controller "tcp:127.0.0.1:6653"
    fail_mode: secure
    Port "s1-eth2"
      Interface "s1-eth2"
    Port "s1-eth1"
      Interface "s1-eth1"
    Port "s1"
      Interface "s1"
        type: internal
    Port "s1-eth3"
      Interface "s1-eth3"
  ovs_version: "2.9.5"
```

Conclusion:

Users follow step-by-step instructions to create a network bridge, connect interfaces, and verify connectivity using ping. This practical approach to learning is fundamental for understanding the creation and management of virtual network infrastructures with OVS, a key technology in Software Defined Networking (SDN).

Open vSwitch Exercise 2: Connect to Controller

The second exercise focuses on connecting an Open vSwitch (OVS) bridge to a reference OpenFlow controller.

Environment:

- **Virtual Machine:** The exercise takes place within a virtual machine likely named "comnetsemu", designed for network emulation.
- **ComNetSEmu Framework:** Provides the foundational tools and environment for network simulation.
- **Mininet:** Used to create the emulated network topology within the VM.
- **Open vSwitch (OVS):** Functions as the virtual switch in this scenario.
- **ovs-testcontroller:** A simple reference OpenFlow controller provided by OVS for testing and demonstration purposes.

Exercise Steps:

1. Secure Fail Mode Configuration:

- Ensure the OVS bridge is set to the "secure" fail mode, which prevents the switch from independently adding flow entries, relying entirely on the controller for flow management.

2. Connect Switch to Controller:

- Establish a connection between the OVS bridge and the reference OpenFlow controller.
- This is likely achieved using a command similar to the one used in the previous exercise, specifying the controller's IP address and port:

```
ovs-vsctl set-controller [bridge name]
tcp:127.0.0.1:6653
```

- In this case, the controller is running locally on port 6653.

```
root@comnetsemu:~/comnetsemu/app/software_defined_networks_examples# ovs-testcontroller -v ptcp:6653
2020-05-06T18:00:45Z|00001|poll_loop|DBG|wakeup due to 0-ms timeout
```

3. Run the Reference Controller:

- In a second terminal window, start the "ovs-testcontroller" to act as the OpenFlow controller: `sudo ovs-testcontroller -v ptcp:6653`
- The `-v` flag increases verbosity, providing more detailed output.
 - `ptcp:6653` specifies the protocol (passive TCP) and the port on which the controller listens.

```
Every 1.0s: ovs-ofctl dump-flows s1                                     comnetsemu: Wed M
NXST_FLOW reply (xid=0x4):
  cookie=0x0, duration=2.255s, table=0, n_packets=0, n_bytes=0, idle_age=2,
CONTROLLER:128
```

4. Dump Flow Entries:

- In another terminal window, use the `ovs-ofctl` utility to display the flow entries currently installed in the OVS bridge: `sudo ovs-ofctl`

```

dump-flows s1
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# ovs-vsctl set-controller
s1 tcp:127.0.0.1:6653
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# ovs-vsctl show
438f8106-c225-4f78-aacc-f2d8724bb3ef
    Bridge "s1"
        Controller "tcp:127.0.0.1:6653"
        fail_mode: secure
        Port "s1-eth2"
            Interface "s1-eth2"
        Port "s1-eth1"
            Interface "s1-eth1"
        Port "s1"
            Interface "s1"
                type: internal
        Port "s1-eth3"
            Interface "s1-eth3"
    ovs_version: "2.9.5"
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# sudo watch -n 1 ovs-ofctl
dump-flows 1
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# sudo watch -n 1 ovs-ofctl
dump-flows s1

```

- This command retrieves and displays the flow table for the bridge named "s1".

Conclusion:

- This exercise demonstrates the practical steps involved in connecting an OVS bridge to an external OpenFlow controller. By setting the switch to "secure" fail mode, we enforce the controller's authority over traffic flow rules.
- Using the `ovs-testcontroller` allows for a simple yet effective way to observe and understand how an OpenFlow controller interacts with an OVS bridge. The `ovs-ofctl dump-flows` command provides visibility into the flow table, which reflects the rules installed by the controller to govern traffic within the emulated network.
- This exercise is fundamental for understanding the concepts and mechanisms of SDN, highlighting the role of the controller as the central point of control in managing network traffic.

Mininet:

Mininet provides a lightweight and scalable platform for creating virtual networks. It leverages Linux container (LXC) technology to create lightweight virtual network elements such as hosts, switches, and routers within a single Linux kernel. These virtual elements can be interconnected to form arbitrary network topologies, facilitating the emulation of various network scenarios.

Creating network Topology

Steps:

1. Install python3-pip

```
rishikesh@Ubuntu:~$ sudo apt install python3-pip
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
python3-pip is already the newest version (22.0.2+dfsg-1ubuntu0.4).
0 upgraded, 0 newly installed, 0 to remove and 216 not upgraded.
rishikesh@Ubuntu:~$
```

2. Install Ryu

```
rishikesh@Ubuntu:~$ sudo pip3 install ryu
[sudo] password for rishikesh:
Requirement already satisfied: ryu in /usr/local/lib/python3.10/dist-packages (4.34)
Requirement already satisfied: six>=1.4.0 in /usr/lib/python3/dist-packages (from ryu) (1.16.0)
Requirement already satisfied: routes in /usr/local/lib/python3.10/dist-packages (from ryu) (2.5.1)
Requirement already satisfied: oslo.config>=2.5.0 in /usr/local/lib/python3.10/dist-packages (from ryu) (9.4.0)
Requirement already satisfied: webob>=1.2 in /usr/local/lib/python3.10/dist-packages (from ryu) (1.8.7)
Requirement already satisfied: netaddr in /usr/local/lib/python3.10/dist-packages (from ryu) (1.2.1)
Requirement already satisfied: tinyrpc in /usr/local/lib/python3.10/dist-packages (from ryu) (1.1.7)
Requirement already satisfied: ovs>=2.6.0 in /usr/lib/python3/dist-packages (from ryu) (2.17.9)
Requirement already satisfied: msgpack>=0.3.0 in /usr/local/lib/python3.10/dist-packages (from ryu) (1.0.8)
Requirement already satisfied: eventlet!=0.18.3,!>=0.20.1,!>=0.21.0,!>=0.23.0,>=0.18.2 in /usr/local/lib/python3.10/dist-packages (from ryu) (0.30.1)
Requirement already satisfied: greenlet>=1.0 in /usr/local/lib/python3.10/dist-packages (from eventlet!=0.18.3,!>=0.20.1,!>=0.21.0,!>=0.23.0,>=0.18.2->ryu) (3.0.3)
Requirement already satisfied: dnspython>=1.15.0 in /usr/local/lib/python3.10/dist-packages (from eventlet!=0.18.3,!>=0.20.1,!>=0.21.0,!>=0.23.0,>=0.18.2->ryu) (2.6.1)
Requirement already satisfied: requests>=2.18.0 in /usr/lib/python3/dist-packages (from oslo.config>=2.5.0->ryu) (2.25.1)
Requirement already satisfied: stevedore>=1.20.0 in /usr/local/lib/python3.10/dist-packages (from oslo.config>=2.5.0->ryu) (5.2.0)
Requirement already satisfied: PyYAML>=5.1 in /usr/lib/python3/dist-packages (from oslo.config>=2.5.0->ryu) (5.4.1)
Requirement already satisfied: oslo.i18n>=3.15.3 in /usr/local/lib/python3.10/dist-packages (from oslo.config>=2.5.0->ryu) (6.3.0)
Requirement already satisfied: debtcollector>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from oslo.config>=2.5.0->ryu) (3.0.0)
Requirement already satisfied: rfc3986>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from oslo.config>=2.5.0->ryu) (2.0.0)
Requirement already satisfied: repoze.lru>=0.3 in /usr/local/lib/python3.10/dist-packages (from routes->ryu) (0.7)
Requirement already satisfied: wrapt>=1.7.0 in /usr/local/lib/python3.10/dist-packages (from debtcollector>=1.2.0->oslo.config>=2.5.0->ryu) (1.16.0)
Requirement already satisfied: pbr!=2.1.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from oslo.i18n>=3.15.3->oslo.config>=2.5.0->ryu) (6.0.0)
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager.
It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv
rishikesh@Ubuntu:~$
```

3. Install git

```
rishikesh@Ubuntu:~$ sudo apt install git
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
git is already the newest version (1:2.34.1-1ubuntu1.10).
0 upgraded, 0 newly installed, 0 to remove and 216 not upgraded.
rishikesh@Ubuntu:~$
```

4. clone git

```
rishikesh@Ubuntu:~$ git clone git://github.com/mininet/mininet
fatal: destination path 'mininet' already exists and is not an empty directory.
rishikesh@Ubuntu:~$ git clone https://github.com/mininet/mininet
fatal: destination path 'mininet' already exists and is not an empty directory.
rishikesh@Ubuntu:~$ ls
Desktop  Downloads  inv.tbh  mealy.v  Music  oflops  openflow  Pictures  Public  Templates  vlsi  vlsi1.zip  xor.cir
Documents  home  inv.v  mininet  myenv  oftest  path  pox  snap  Videos  vlsi1  vlsi.zip
rishikesh@Ubuntu:~$
```

5. Command promote environment

```
rishikesh@Ubuntu:~$ cd mininet/
rishikesh@Ubuntu:~/mininet$ cd util/
rishikesh@Ubuntu:~/mininet/util$ ls
build-ovs-packages.sh  colorfilters  install.sh  m          openflow-patches  sysctl-addon  versioncheck.py
clustersetup.sh         doxify.py     kbuild    nox-patches  sch_htb-ofbuf   unpep8      VM
rishikesh@Ubuntu:~/mininet/util$ sudo./install.sh -a
bash: sudo./install.sh: No such file or directory
```

6. execute script named "install.sh" with administrative privileges

```
rishikesh@Ubuntu:~/mininet/util$ sudo ./install.sh -a
Detected Linux distribution: Ubuntu 22.04 jammy amd64
sys.version_info(major=3, minor=10, micro=12, releaselevel='final', serial=0)
Detected Python (python3) version 3
Installing all packages except for -eix (doxypy, ivs, nox-classic)...
Install Mininet-compatible kernel if necessary
Get:1 http://security.ubuntu.com/ubuntu jammy-security InRelease [110 kB]
Hit:2 http://in.archive.ubuntu.com/ubuntu jammy InRelease
Get:3 http://in.archive.ubuntu.com/ubuntu jammy-updates InRelease [119 kB]
Hit:4 http://in.archive.ubuntu.com/ubuntu jammy-backports InRelease
Fetched 229 kB in 3s (65.7 kB/s)
Reading package lists... Done
Reading package lists...
Building dependency tree...
Reading state information...
linux-image-6.2.0-39-generic is already the newest version (6.2.0-39.40~22.04.1).
0 upgraded, 0 newly installed, 0 to remove and 216 not upgraded.
Installing Mininet dependencies
Reading package lists...
Building dependency tree...
Reading state information...
gcc is already the newest version (4:11.2.0-1ubuntu1).
make is already the newest version (4.3-4.1build1).
net-tools is already the newest version (1.60+git20181103.0eebece-1ubuntu5).
psmisc is already the newest version (23.4-2build3).
socat is already the newest version (1.7.4.1-3ubuntu4).
telnet is already the newest version (0.17-44build1).
help2man is already the newest version (1.49.1).
iperf is already the newest version (2.1.5+dfsg1-1).
pep8 is already the newest version (1.7.1-9ubuntu1).
pyflakes3 is already the newest version (2.4.0-2).
```

```
/home/rishikesh/mininet/util
Installing OpenFlow reference implementation...
Reading package lists...
Building dependency tree...
Reading state information...
autoconf is already the newest version (2.71-2).
automake is already the newest version (1:1.16.5-1.3).
gcc is already the newest version (4:11.2.0-1ubuntu1).
libtool is already the newest version (2.4.6-15build2).
make is already the newest version (4.3-4.1build1).
0 upgraded, 0 newly installed, 0 to remove and 216 not upgraded.
Reading package lists...
Building dependency tree...
Reading state information...
autotools-dev is already the newest version (20220109.1).
pkg-config is already the newest version (0.29.2-1ubuntu3).
git is already the newest version (1:2.34.1-1ubuntu1.10).
libc6-dev is already the newest version (2.35-0ubuntu3.7).
0 upgraded, 0 newly installed, 0 to remove and 216 not upgraded.
fatal: destination path 'openflow' already exists and is not an empty directory.
rishikesh@Ubuntu:~/mininet/util$
```

7. Running Mininet with elevated privileges for creating and testing network configurations requiring administrative access.

```
rishikesh@Ubuntu:~/mininet/util$ sudo mn
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1)
*** Configuring hosts
h1 h2
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2
h2 -> h1
*** Results: 0% dropped (2/2 received)
mininet> exit
*** Stopping 1 controllers
c0
*** Stopping 2 links
..
*** Stopping 1 switches
s1
*** Stopping 2 hosts
h1 h2
*** Done
completed in 29.512 seconds
```

8. Creating a network with one switch and three hosts, configuring MAC addresses, using the Open vSwitch, and setting a remote controller using sudo privileges.

```
rishikesh@Ubuntu:~/mininet/util$ sudo mn --topo single,3 --mac --switch ovsk --controller remote
*** Creating network
*** Adding controller
Unable to contact the remote controller at 127.0.0.1:6653
Unable to contact the remote controller at 127.0.0.1:6633
Setting remote controller to 127.0.0.1:6653
*** Adding hosts:
h1 h2 h3
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1) (h3, s1)
*** Configuring hosts
h1 h2 h3
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet> pingall
*** Ping: testing ping reachability
h1 -> X X
h2 -> X X
h3 -> X X
*** Results: 100% dropped (0/6 received)
mininet> exit
*** Stopping 1 controllers
c0
*** Stopping 3 links
...
*** Stopping 1 switches
s1
*** Stopping 3 hosts
h1 h2 h3
*** Done
completed in 43.889 seconds
rishikesh@Ubuntu:~/mininet/util$
```

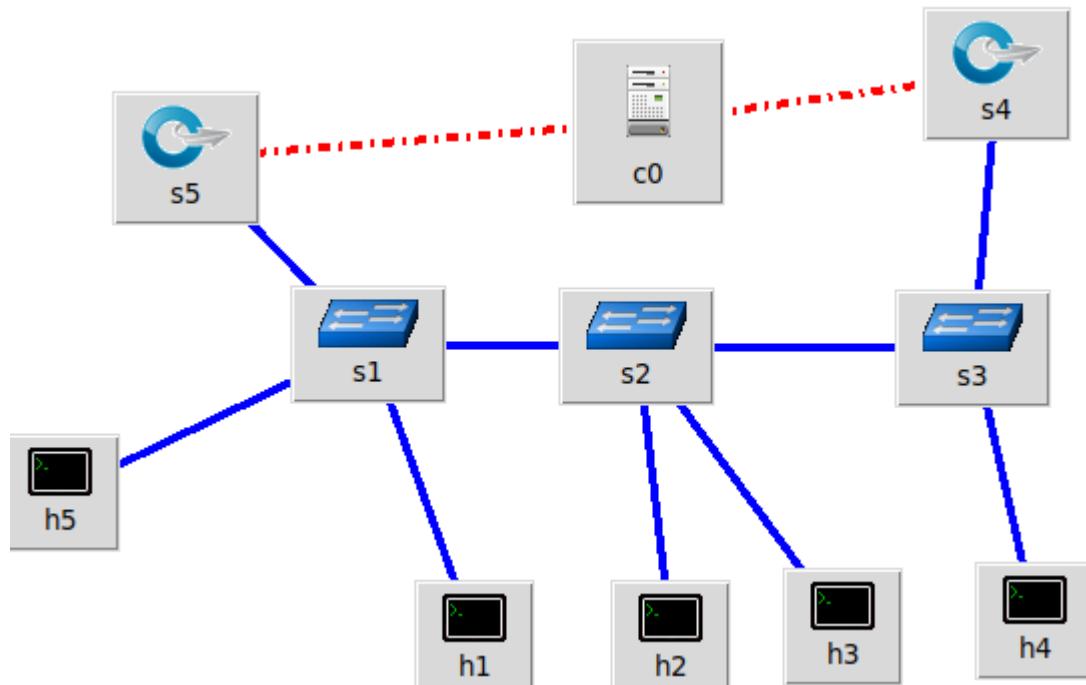
9.Launching the Mininet GUI editor from the specified directory with sudo privileges.

```
rishikesh@Ubuntu:~$ sudo ~/mininet/examples/miniedit.py
[sudo] password for rishikesh:
/usr/bin/env: 'python': No such file or directory
rishikesh@Ubuntu:~$ sudo apt install python3
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
python3 is already the newest version (3.10.6-1~22.04).
0 upgraded, 0 newly installed, 0 to remove and 216 not upgraded.
rishikesh@Ubuntu:~$
```

```
rishikesh@Ubuntu:~$ cd mininet/
rishikesh@Ubuntu:~/mininet$ ls
bin      custom  examples  Makefile      mn.1      mnexec.c  util
build    debian  INSTALL   mininet     mnexec    README.md
CONTRIBUTORS  doc    LICENSE   mininet.egg-info mnexec.1  setup.py
rishikesh@Ubuntu:~/mininet$ cd examples/
```

10. Executing the Mininet GUI editor script using Python 3 with administrative privileges.

```
rishikesh@Ubuntu:~/mininet/examples$ sudo python3 ./miniedit.py
topo=None
```



3. RyU SN Controller: Example: Sending a message from switch

Ryu Applications :

- my_app.py: from ryu.base import app_manager from ryu.ofproto import ofproto_v1_3
class MyApp(app_manager.RyuApp): OFP_VERSIONS = [ofproto_v1_3.OFP_VERSION]
- Run application with: ryu-manager my_app.py ss

Asynchronous Messages :

- Messages from switches triggered by events
- Call function inside the application class.
- Example: @set_ev_cls(ofp_event.EventOFPPacketIn, MAIN_DISPATCHER) def packet_in_handler(self, ev):
msg= ev.msg dp = msg.datapath ofp = dp.ofproto ofp_parser =
dp.ofproto_parser

Controller to Switch Messages:

- Messages sent from the controller to instruct the switch
- Eg. forward packets, add flows, query statistics, ..
- Example: flood a packet
actions = [ofp_parser.OFPActionOutput(ofp.OFPP_FLOOD)] out =
ofp_parser.OFPPacketOut(datapath=dp, buffer_id=msg.buffer_id,
in_port=msg.in_port, actions=actions) dp.send_msg(out)

Exercise : 3 – Ryu Examples

Run the Ryu application hub.py (stored in ~/sdn/ex03/hub.py)

ryu-manager --verbose hub.py

- Open second terminal
- Run mininet and connect to the controller

- sudo mn --topo single,3 \ --controller remote
 - Send some packets with mininet, e.g. with „pingall“ or „iperf“
 - Observe the output in Ryu
 - Repeat with switch_13.py and sdn_switch_13.py
- ryu-manager --verbose hub.py

Install Ryu and mininet on Ubuntu :

```
sudo apt update  
sudo apt install ryu  
ryu-manager --version  
sudo apt update  
git clone git://github.com/mininet/mininet  
cd mininet  
sudo ./util/install.sh -a  
sudo mn --test pingall
```

After using all commands given up you can connect to the github and install the require commands like ryu and mininet

```

Terminal
File Edit View Search Terminal Tabs Help
Battery low
Approximately 58 minutes
eswar@eswar-VirtualBox: ~$ sudo apt update
[sudo] password for eswar:
Get:1 http://in.archive.ubuntu.com/ubuntu jammy InRelease [110 kB]
Get:2 http://in.archive.ubuntu.com/ubuntu jammy-updates InRelease [110 kB]
Get:3 http://in.archive.ubuntu.com/ubuntu jammy-backports InRelease [110 kB]
Get:4 http://security.ubuntu.com/ubuntu jammy-security InRelease [110 kB]
Get:5 http://security.ubuntu.com/ubuntu jammy-security/main amd64 Packages [1,562 kB]
Get:6 http://security.ubuntu.com/ubuntu jammy-security/main i386 Packages [110 kB]
Get:7 http://security.ubuntu.com/ubuntu jammy-security/main amd64 Packages [1,346 kB]
Get:8 http://security.ubuntu.com/ubuntu jammy-security/main i386 Packages [404 kB]
Get:9 http://security.ubuntu.com/ubuntu jammy-security/universe amd64 Packages [110 kB]
Get:10 http://security.ubuntu.com/ubuntu jammy-security/universe i386 Packages [660 kB]
Get:11 http://in.archive.ubuntu.com/ubuntu jammy-updates/main Translation-en [128 kB]
Get:12 http://in.archive.ubuntu.com/ubuntu jammy-updates/universe amd64 Packages [110 kB]
Get:13 http://in.archive.ubuntu.com/ubuntu jammy-updates/universe i386 Packages [110 kB]
Get:14 http://in.archive.ubuntu.com/ubuntu jammy-updates/restricted amd64 Packages [1,754 kB]
Get:15 http://in.archive.ubuntu.com/ubuntu jammy-updates/restricted i386 Packages [38.1 kB]
Get:16 http://in.archive.ubuntu.com/ubuntu jammy-updates/multiverse amd64 Packages [110 kB]
Get:17 http://in.archive.ubuntu.com/ubuntu jammy-updates/multiverse i386 Packages [110 kB]
Get:18 http://in.archive.ubuntu.com/ubuntu jammy-updates/multiverse amd64 Packages [42.7 kB]
Get:19 http://in.archive.ubuntu.com/ubuntu jammy-updates/multiverse i386 Packages [4,472 kB]
Get:20 http://in.archive.ubuntu.com/ubuntu jammy-updates/multiverse Translation-en [10.4 kB]
Fetched 1,084 MB in 10s (108 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
271 packages and 0 global ones installed. Run 'apt list --upgradable' to see them.
eswar@eswar-VirtualBox: ~$ sudo mn --test pingall
sudo: mn: command not found
eswar@eswar-VirtualBox: ~$ sudo mn --test pingall
sudo: mn: command not found
eswar@eswar-VirtualBox: ~$ git clone git://github.com/mininet/mininet
Cloning into 'mininet'...
fatal: unable to connect to github.com:
github.com[0: 20.207.73.82]: errno=Connection timed out
eswar@eswar-VirtualBox: ~$ git clone https://github.com/mininet/mininet.git
Cloning into 'mininet'...
remote: Enumerating objects: 10388, done.
remote: Counting objects: 100% (234/234), done.
remote: Compressing objects: 100% (139/139), done.
remote: Total 10388 (delta 93), reused 176 (delta 93), pack-reused 10154

```

Next use the command

- Sudo mn-- topo single, 3 --controller remote,ip
for creating a network and adding controllers and then to test pingall

```

Activities Terminal
File Edit View Search Terminal Tabs Help
File Edit View Search Terminal Tabs Help
eswar@eswar-VirtualBox: ~/mininet
*** Creating network
*** Adding controller
Unable to contact the remote controller at 127.0.0.1:6633
*** Adding hosts:
h1 h2 h3
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1) (h3, s1)
*** Configuring hosts
h1 h2 h3
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet> pingall
*** Ping: testing ping reachability
h1 -> X X
h2 -> X X
h3 -> X X
*** Results: 100% dropped (0/6 received)
mininet> s1 xterm &
mininet>

```

On testing the pingall with some values like (h1 h2) it will open another tab of root where we can see the NXST_FLOW reply (xid=0*4)

```
root@eswar-VirtualBox: /home/eswar/mininet - □ ×
Every 1.0s: ovs-ofctl dump-flows s1    eswar-VirtualBox: Thu Apr 18 15:24:50 2024 [Bo]
* C|NXST_FLOW reply (xid=0x4):
* A|
able|
* A|
```

You can check your input messages values here

Like you want mininet>h1 ping h2

```
mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=6.52 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=6.17 ms
```

```
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# sudo watch
-n 1 ovs-ofctl dump-flows s1
root@comnetsemu:/comnetsemu/app/software_defined_networks_examples# tcpdump -i
s1-eth3
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on s1-eth3, link-type EN10MB (Ethernet), capture size 262144 bytes
18:42:09.649753 IP 10.0.0.1 > 10.0.0.2: ICMP echo request. id 4790, seq 1, length 64
18:42:09.654396 IP 10.0.0.2 > 10.0.0.1: ICMP echo reply. id 4790, seq 1, length 64
18:42:10.650515 IP 10.0.0.1 > 10.0.0.2: ICMP echo request. id 4790, seq 2, length 64
18:42:10.655416 IP 10.0.0.2 > 10.0.0.1: ICMP echo reply. id 4790, seq 2, length 64
```

You can see the packages loading

Now in the second terminal we are going to run a command nano hub.py

```

# here we gather the actions the switch should execute with the packet
# OFPActionOutput instructs the switch to send the message to an output port
# OFPP_FLOOD is a special output port for sending on all output ports
actions = [ofp_parser.OFPActionOutput(ofproto.OFPP_FLOOD)]

# if the packet is not in the switch buffer, include it in the PacketOut message
data = None
if msg.buffer_id == ofproto.OFP_NO_BUFFER:
    data = msg.data

# we put the actions into a packet_out message
# the message also contains
out = ofp_parser.OFPPacketOut(
    datapath=dp, # assign the message to the switch
    buffer_id=msg.buffer_id, # reflect the id assigned by the switch
    in_port=in_port, # input port
    actions=actions, # our list of actions
    data=data) # send the packet back

```

Here we can see all the packages used and upcoming packages

Now by using ryu-manager sdn_switch_13.py --verbose command So, when you run `ryu-manager sdn_switch_13.py --verbose`, you're essentially starting the Ryu controller with the specified application (`sdn_switch_13.py`) and opting to receive verbose output, which can aid in troubleshooting and monitoring the SDN environment.

After do the ping test for mininet>h1 ping h2

```

mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=14.1 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=5.59 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.414 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.132 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.144 ms
64 bytes from 10.0.0.2: icmp_seq=6 ttl=64 time=0.069 ms
64 bytes from 10.0.0.2: icmp_seq=7 ttl=64 time=0.120 ms
64 bytes from 10.0.0.2: icmp_seq=8 ttl=64 time=0.143 ms

```

```
NXST_FLOW reply (xid=0x4):
cookie=0x0, duration=38.317s, table=0, n_packets=13, n_bytes=1162, idle_age=27, in_port=2,dl_dst=2e:3a:c5:75:90:47 actions=output:1
cookie=0x0, duration=37.325s, table=0, n_packets=13, n_bytes=1162, idle_age=27, in_port=1,dl_dst=22:ae:5d:c7:49:57 actions=output:2
cookie=0x0, duration=44.352s, table=0, n_packets=4214, n_bytes=6443668, idle_age=38, priority=0 actions=CONTROLLER:65535
```

As you can see the difference here it is more simple and fast

NXST FLOW reply: This might indicate the software or tool used to generate the information, possibly related to network flow analysis.

xid: This stands for transaction ID, a unique identifier for this specific network flow.

cookie: This value is likely used for internal bookkeeping by the network analysis tool.

duration: This represents the amount of time (in seconds) that this network flow was active.

table: This could refer to a routing table or table used for managing network flows.

n_packets: This indicates the total number of data packets included in the network flow.

n_bytes: This represents the total size (in bytes) of all the data packets in the network flow.

idle_age: This might indicate how long it's been since the last packet was received in this flow.

in_port: This refers to the network port that the data packets in this flow originated from.

dl_dst: This stands for "destination link layer address," which is the MAC address of the device that received the data packets.

actions: This section likely specifies what actions were taken when processing this network flow (e.g., routing it to a specific output port).

Next use the command

```
mininet>pingall
```

Then use mininet>iperf for testing bandwidth between h1 and h3

As You can see the result for that its 16.1 Gbits/sec

```

mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3
h2 -> h1 h3
h3 -> h1 h2
*** Results: 0% dropped (6/6 received)
mininet> iperf
*** Iperf: testing TCP bandwidth between h1 and h3
*** Results: ['16.1 Gbits/sec', '16.1 Gbits/sec']
mininet> █

```

```

NXST_FLOW reply (xid=0x4):
cookie=0x0, duration=150.686s, table=0, n_packets=17, n_bytes=1442, idle_age=41, in_port=2,dl_dst=2e:3a:c5:75:90:47 actions=output:1
cookie=0x0, duration=149.694s, table=0, n_packets=17, n_bytes=1442, idle_age=41, in_port=1,dl_dst=22:ae:5d:c7:49:57 actions=output:2
cookie=0x0, duration=46.897s, table=0, n_packets=82851, n_bytes=5468886, idle_age=2, in_port=3,dl_dst=2e:3a:c5:75:90:47 actions=output:1
cookie=0x0, duration=46.872s, table=0, n_packets=3, n_bytes=182, idle_age=41, in_port=2,dl_dst=da:f2:f7:06:71:b0 actions=output:3
cookie=0x0, duration=46.869s, table=0, n_packets=3, n_bytes=182, idle_age=41, in_port=3,dl_dst=22:ae:5d:c7:49:57 actions=output:2
cookie=0x0, duration=46.864s, table=0, n_packets=195105, n_bytes=10088577802, idle_age=2, in_port=1,dl_dst=da:f2:f7:06:71:b0 actions=output:3
cookie=0x0, duration=156.721s, table=0, n_packets=4218, n_bytes=6444060, idle_age=46, priority=0 actions=CONTROLLER:65535

```

As we can see that the duration is increasing and package

```

if dst in self.mac_to_port[dpid]:
    # if we already know the mac we can select the output port
    out_port = self.mac_to_port[dpid][dst]
    self.logger.info('Message for known {} at port {}'.format(src, out_port))

    # if we know the destination we can establish a flow
    self.add_flow(datapath, dst, in_port, out_port)
else:
    # but if it can't be found we just flood the packet
    self.logger.info('Message for unknown {}'.format(src))
    out_port = ofproto.OFPP_FLOOD

```

function used for processing data packets

```

def add_flow(self, datapath, dst, in_port, out_port):
    ofproto = datapath.ofproto
    ofp_parser = datapath.ofproto_parser

    # to add a flow we have to define some match criteria
    match = ofp_parser.OFPMatch(
        in_port=in_port, # we match the input port
        eth_dst=dst) # and the destination MAC address

    # we also need to define the desired action
    actions = [ofp_parser.OFPActionOutput(out_port)]
    inst = [ofp_parser.OFPIInstructionActions(ofproto.OFFIT_APPLY_ACTIONS,
                                              actions)]
    # now we build a flow table entry
    flow = ofp_parser.OFPFlowMod(
        datapath=datapath, # assign datapath

```

Function Definition: This code defines a function named `add_flow` that takes four arguments:

- `datapath`: This represents the switch or network device where the flow entry will be added.
- `dst`: This is the destination MAC address that the flow entry should match.
- `in_port`: This is the input port number that the flow entry should match.
- `out_port`: This is the output port number to which packets matching the flow entry should be sent.

Match Criteria:

- The code defines a match object using `ofp_parser.OFPMatch`. This object specifies the criteria that packets must meet in order for the flow entry to be applied.
- In this case, the match criteria include:
 - `in_port`: Packets must arrive on the specified input port (`in_port`).
 - `eth_dst`: Packets must have a destination MAC address that matches the specified value (`dst`).

Actions:

- The code defines a list of actions using `ofp_parser.OFPActionOutput`. This list specifies what actions should be taken when a packet matches the flow entry.
- In this case, the only action is to send the packet to the specified output port (`out_port`).

Instruction:

- The code combines the match criteria and actions into an instruction using `ofp_parser.OFPInstructionActions`. This instruction tells the switch what to do with packets that match the flow entry.

Flow Entry:

- The code creates a flow entry object using `ofp_parser.OFFPFlowMod`. This object encapsulates all the information about the flow, including:
 - `datapath`: The switch where the flow entry will be added (same as the function argument).
 - `match`: The matching criteria defined earlier.
 - `instructions`: The instructions to be applied to matching packets.
 - `priority` (optional): An arbitrary priority level for the flow entry (higher priority entries are more likely to be applied). In this case, a priority of 1000 is set.

Adding the Flow:

- Finally, the code sends the newly created flow entry (`flow`) to the switch using `datapath.send_msg`. This adds the flow entry to the flow table on the switch.

More SDN Applications

- Smarter firewalls (distributed, content sensitive, ...)
- Full network load balancing (Multipath)
- Quality of Service (metering supported since OpenFlow 1.3)
- Network slicing
- Intrusion detection systems
- Overlay networks (e.g. VPN)
- Resilience and Failover mechanisms